

Some practical considerations

Using our own data

- suppose we want to build an arabic digit classification system



- data is give to us in, say, two spreadsheets
 - images spreadsheet with columns containing uint8 data
 - labels spreadsheet containing a single column of int data

Inheriting from the Dataset class

```
import pandas as pd
from torch.utils.data import Dataset

class AMNIST(Dataset):
    def __init__(self, imagesfile, labelsfile):
        super().__init__()
        df_x = pd.read_csv(imagesfile, header=None)
        df_y = pd.read_csv(labelsfile, header=None)
        self.X = df_x.iloc[:, :].to_numpy().reshape(-1,28,28) / 255.0
        self.y = df_y.iloc[:, 0].to_numpy()

    def __len__(self):
        return len(self.y)

    def __getitem__(self, idx):
        image = self.X[idx]
        image = torch.tensor(np.expand_dims(image, axis=0))
        label = torch.tensor(self.y[idx])
        return image, label
```

Saving and loading models

```
# saving/loading models  
PATH = './cifar_net.pth'  
torch.save(net.state_dict(), PATH)  
  
net = Net()  
net.load_state_dict(torch.load(PATH))
```

Example of finetuning an existing network

```
from torchvision import models
```

```
net = models.resnet18(pretrained=True)
```


```
for layer in net.parameters():
```

```
    layer.requires_grad=False
```

```
net.conv1 = nn.Conv2d(1, 64, kernel_size=(7, 7), stride=(2, 2),
```

```
    ↪ padding=(3, 3), bias=False)
```

```
net.fc = nn.Linear (in_features=512, out_features=10, bias=True)
```



Data transformations

```
1 from torchvision import transforms
2
3 transform_train = transforms.Compose([
4     transforms.RandomCrop(28, padding=4),
5     transforms.ToTensor(),
6 ])
7 # Training dataset
8 train_data = MNIST(root='./datasets', train=True, download=True,
9     ↪ transform=transform_train)
10
11 train_loader = torch.utils.data.DataLoader(train_data,
12     ↪ batch_size=batch_sz, shuffle=True, pin_memory=True)
```

Guess what this does:

```
1 dataset = MyMNIST(image_dir, label_spreadsheet,
2     ↪ transforms.Compose([transforms.ToTensor(),
3     ↪ transforms.ColorJitter(...), transforms.RandomAffine(90)]))
4 train_loader=DataLoader(dataset, batch_size=10)
```

Standard deviation and data normalization

- $\text{avg}(x) = \mathbf{1}^T x / n$
- de-meaned vector is $\tilde{x} = x - \text{avg}(x)\mathbf{1}$ (so $\text{avg}(\tilde{x}) = 0$)
- standard deviation of x

$$\text{std}(x) = \text{rms}(\tilde{x}) = \frac{\|x - (\mathbf{1}^T x / n)\mathbf{1}\|}{\sqrt{n}}$$

- $\text{std}(x)$ gives “typical amount” x_i vary from $\text{avg}(x)$
- $\text{std}(x) = 0$ only if $x = \alpha\mathbf{1}$ for some α
- a basic formula

$$\text{rms}(x)^2 = \text{avg}(x)^2 + \text{std}(x)^2$$

- standardization ($\mu = 0$, $\sigma = 1$, z-scores)

$$z = \frac{1}{\text{std}(x)}(x - \text{avg}(x)\mathbf{1})$$

Three sources of errors

- error between a trained neural network and a perfect classification function to learn can be viewed as having three components
 - optimization error
 - approximation error
 - generalization error

Practical training considerations



- the exploding gradient problem—gradient clipping solution

$$g' = \min(1, \frac{c}{\|g\|})g$$


- the vanishing gradient problem
 - non-saturating activation functions
 - residual networks (ResNets)
 - parameter initialization
 - standardize activations at each layer

Non-saturating activation functions

Name	Definition
Sigmoid	$\sigma(a) = \frac{1}{1+e^{-a}}$
Hyperbolic tangent	$\tanh(a) = 2\sigma(2a) - 1$
Softplus	$\sigma_+(a) = \log(1 + e^a)$
Rectified linear unit	$\text{ReLU}(a) = \max(a, 0)$
Leaky ReLU	$\max(a, 0) + \alpha \min(a, 0)$
Exponential linear unit	$\max(a, 0) + \min(\alpha(e^a - 1), 0)$
Swish	$a\sigma(a)$
GELU	$a\Phi(a)$



ResNets



- intuition is that it is easier to learn to generate a small perturbation to the input than to directly predict the output
- gradient at layer l depends on the gradient of layer L in a way that is independent of the depth of the network

Parameter initialization

- initialization of θ parameters need to be done somewhat carefully
- sampling parameters from a standard normal with fixed variance can result in exploding activations or gradients
- pytorch has a (good) Glorot initialization, $\sigma^2 = 2/(n_{in} + n_{out})$
- data-driven initializations, which compute variances of the activations across a minibatch, may be used



Regularization

- it is common to add a regularization term to the NLL loss objective



$$\mathcal{L}(\theta) + \text{regularization}$$

- $\lambda \|\theta\|^2$ often used for regularization
 - encourage small weights, and therefore simpler models
 - referred to as weight decay
- stopping when the error on validation starts to increase is a regularization strategy that prevents overfitting
 - referred to as early stopping
 - but... double descent curves have been often observed
- dropout can also prevent overfitting
- encourage flat minima
- and many others....

Double descent in overparameterized networks




Dropout



- turn off (randomly, on a per-example, or per-batch basis) all outgoing connections from each neuron with probability p
- intuitively, each neuron must learn to perform well even if others units are randomly missing—prevents complex and fragile dependencies from being learned

Local minima are not all created equal



- flat minima correspond to regions in parameters space where there is a lot of posterior uncertainty
- hence solutions from this region are unlikely to memorize irrelevant details from the training set
- noise in SGD may prevent algorithm from entering narrow regions of landscape
- sharpness aware terms can be added the objective
- estimates of variance of minibatch gradients (a measure of stability, and hence generalization ability) may be added

Batch normalization layer

- in order to reduce vanishing or exploding gradient problems with SGD, it is often useful to normalize data in every layer
- similar to standardizing input data
- batch normalization: replace activation vector z_n by \tilde{z}_n

$$\mu_{\mathcal{B}} = \frac{1}{|\mathcal{B}|} \sum_{z \in \mathcal{B}} z$$


$$\sigma_{\mathcal{B}}^2 = \frac{1}{|\mathcal{B}|} \sum_{z \in \mathcal{B}} (z - \mu_{\mathcal{B}})^2$$

$$\hat{z}_n = \frac{(z - \mu_{\mathcal{B}})}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}}$$

$$\tilde{z}_n = \gamma \hat{z}_n + \beta$$

- done dynamically in every layer with every batch
 - trainable parameters: one γ and one β per feature map per layer
 - note: must be held fixed at inference time (post training)

Other kinds of normalizations



- can be conveniently added as layers in the neural network definition

Optimization for training NN



- learn a function $f(x_n; \theta) \approx y_n$
- loss $\mathcal{L}(\theta) = \frac{1}{N} \sum_{n=1}^N \ell(y, f(x_n; \theta))$
 - “loss” incurred for not properly aligning our prediction $f(x)$ with y
- regularization to avoid overfitting

$$\theta^* = \operatorname{argmin}_{\theta} \mathcal{L}(\theta) + \lambda \operatorname{pen}(\theta)$$

- $\operatorname{pen}(\theta)$ takes lower values for parameters θ that yield functions f with lower complexity, λ is a regularization parameter that can be tuned
- local descent, compute next iterate as $\theta^{(k+1)} = \theta^{(k)} + \alpha^{(k)} d^{(k)}$
- choice of learning rate $\alpha^{(k)}$
 - fixed α called learning rate
 - decaying step factor

$$\alpha^{(k)} = \alpha^{(1)} \gamma^{k-1}$$

Decaying learning rate



- fixed learning rate $\alpha^{(k)} = \alpha$
- vs exponentially decaying learning rate $\alpha^{(k)} = \alpha 0.99^k$

Stochastic gradient descent

- search in direction of steepest descent (negative gradient direction)

$$d^{(k)} = -\nabla f(\theta^{(k)})$$

- in our use case of interest, our objective is a loss of the form:

$$\frac{1}{N} \sum_{n=1}^N \ell(y_n, f(x_n; \theta))$$

- a noisy/randomized gradient may be obtained by choosing a subset of the training data (batch)

Termination conditions

- maximum iterations

$$k > k_{\max}$$

- absolute improvement

$$f(\theta^{(k)}) - f(\theta^{(k+1)}) < \epsilon_a$$


- relative improvement

$$f(\theta^{(k)}) - f(\theta^{(k+1)}) < \epsilon_r |f(\theta^{(k)})|$$

- gradient magnitude


$$\|\nabla f(\theta^{(k+1)})\| \leq \epsilon_g$$

Momentum



Momentum

- some functions cause gradient descent to get stuck
- momentum overcomes these issues by replicating the effect of physical momentum



Further improvements

- Adagrad
 - instead of using the same learning rate for all components of x , Adaptive Subgradient method (Adagrad) adapts the learning rate for each component of x
 - dulls the influence of parameters with consistently high gradients
- RMSprop
 - extends Adagrad to avoid monotonically decreasing learning rate by maintaining a decaying average of squared gradients
- Adam
 - incorporates ideas from momentum and RMSProp
 - the adaptive moment estimation method (Adam), adapts the learning rate to each parameter