

# **ADVANCED ARTIFICIAL INTELLIGENCE: COMPUTER VISION**

(tips and hints for lab work)

## Attendance




## Today's labs

- Lab: add batch normalization to CNN for CIFAR10
- Lab: finetune ResNet18 with CIFAR10 (or MNIST) data
  - retrain the whole network from the pretrained weights
  - freeze part of the network (at the pretrained weights) and train the rest
- Lab: generate MNIST-like images
- Lab: generate CIFAR-like images

## Note on image data in custom data loader

- jpg images are 3-channel color images (we see them as grayscale, when R, G, and B channels have the same values)
- image data values stored as `uint8` (values between 0 and 255)
- for display image data must be stored in the format  $H \times W \times C$
- but pytorch requires data to be in  $C \times H \times W$  and `float32` (typically in the range 0.0–1.0)
- transform from `uint8 HWC` to `float CHW` either:
  - explicitly in the `__getitem__` method, or
  - by using `ToTensor()` transformation  
(`transform=transforms.ToTensor()`)

## Batch normalization layer



## Batch normalization layer I

- in order to reduce vanishing or exploding gradient problems with SGD, it is often useful to normalize data in every layer
- similar to standardizing input data
- batch normalization: replace activation vector  $z_n$  by  $\tilde{z}_n$

$$\mu_{\mathcal{B}} = \frac{1}{|\mathcal{B}|} \sum_{z \in \mathcal{B}} z$$


$$\sigma_{\mathcal{B}}^2 = \frac{1}{|\mathcal{B}|} \sum_{z \in \mathcal{B}} (z - \mu_{\mathcal{B}})^2$$

$$\hat{z}_n = \frac{(z - \mu_{\mathcal{B}})}{\sqrt{\sigma_{\mathcal{B}}^2 + \epsilon}}$$

$$\tilde{z}_n = \gamma \hat{z}_n + \beta$$

## Batch normalization layer II

- done dynamically in every layer with every batch
  - trainable parameters: one  $\gamma$  and one  $\beta$  per feature map per layer
  - note: mean and variance should be held fixed at inference time (`net.eval()`, `net.train()`)




```
1  class CNN_bn(nn.Module):
2      def __init__(self):
3          super().__init__()
4          self.conv1 = nn.Conv2d(3, 6, 5)
5          self.pool = nn.MaxPool2d(2, 2)
6          self.bn1 = nn.BatchNorm2d(6)
7
8          self.conv2 = nn.Conv2d(6, 16, 5)
9          self.bn2 = nn.BatchNorm2d(16)
10
11         self.fc1 = nn.Linear(16 * 5 * 5, 120)
12         self.bn3 = nn.BatchNorm1d(120)
13
14         self.fc2 = nn.Linear(120, 84)
15         self.bn4 = nn.BatchNorm1d(84)
16
17         self.fc3 = nn.Linear(84, 10)
18
19     def forward(self, x):
20         x = self.bn1(self.pool(F.relu(self.conv1(x))))
21         x = self.bn2(self.pool(F.relu(self.conv2(x))))
22         x = torch.flatten(x, 1)
23         x = F.relu(self.fc1(x))
24         x = F.relu(self.fc2(x))
25         x = self.fc3(x)
26         return x
```



## Two modes for forward pass: training vs evaluation


```
1 net.eval()
2
3 correct = 0
4 total = 0
5 # no need to compute gradients, since we are not training
6 with torch.no_grad():
7     for data in test_loader:
8         images, labels = data
9
10        # run images through the network
11        outputs = net(images)
12
13        # class with largest value is our prediction
14        _, predicted = torch.max(outputs.data, 1)
15
16        total += labels.size(0)
17        correct += (predicted == labels).sum().item()
18
19 print(f'Accuracy on test images: {100 * correct // total} %')
```

# Transfer learning



# Finetuning Resnet18 — starting from a pre-trained network



```
1 from torchvision.models import resnet18, ResNet18_Weights
2 weights = ResNet18_Weights.DEFAULT
3 model = resnet18(weights=weights, progress=False)
4
5 input_features = model.fc.in_features
6 model.fc = nn.Linear(in_features=input_features, out_features=10,
   ↪ bias=True)
7
8 model.conv1 = nn.Conv2d(3, 64, kernel_size=(3, 3), stride=1, padding=1,
   ↪ bias=False)
```



## Finetuning Resnet18 — Freezing layers

```
1  for param in model.parameters():  
2      param.requires_grad = False  
3  
4      # model.conv1.requires_grad_(True)  
5  model.layer4.requires_grad_(True)  
6  model.fc.requires_grad_(True)
```

## Strided convolution



- with a filter of size  $f_h \times f_w$ , an image of size  $x_h \times x_w$ , and padding of size  $p_h$  and  $p_w$  on each side, the output size is

$$(x_h + 2p_h - f_h + 1) \times (x_w + 2p_w - f_w + 1)$$

- if we perform the convolution every  $s$  pixels, output size is

$$\left( \left\lfloor \frac{x_h + 2p_h - f_h}{s_h} \right\rfloor + 1 \right) \times \left( \left\lfloor \frac{x_w + 2p_w - f_w}{s_w} \right\rfloor + 1 \right)$$